

DEVELOPER ROADMAP

Stop Paying for GitHub Copilot: The 5-Step Local AI Setup

A complete step-by-step setup guide for a private, offline developer loop.

SLIDE 2

Step 1: Sizing to Your Workstation

The VRAM Bottleneck

If parameters exceed GPU VRAM, the system swaps to RAM. Generation drops to **< 2 tokens/sec**. Keep models strictly inside VRAM.

Rule of Thumb: Size model footprint to fit within 70% of total VRAM, leaving 30% for OS/browser.

Hardware Compatibility

- **8 GB VRAM:** < 3B models (e.g. Qwen 1.5B Base)
- **16-24 GB VRAM:** 12B-14B models (e.g. Qwen 14B / Gemma 12B)
- **32 GB+ VRAM:** 22B+ models (e.g. Codestral 22B)

SLIDE 3

Step 2: Install Engine & Pull Weights



```
# 1. Download & Install Ollama from ollama.com
```

```
# 2. Pull base model for tab autocomplete  
$ ollama pull qwen2.5-coder:1.5b-base
```

```
# 3. Pull instruct model for chat reasoning  
$ ollama pull qwen2.5-coder:14b-instruct
```

Ollama schedules GPU memory and runs the local API server on port 11434.

SLIDE 4

Step 3: IDE Split-Model Configuration

1. config.json Configuration

```
{
  "models": [{"provider":
"ollama", "model": "qwen2.5-
coder:14b"}],
  "tabAutocompleteModel": {
    "provider": "ollama",
    "model": "qwen2.5-coder:1.5b-
base"
  }
}
```

Replacing Copilot Features

- **Ghost Text**

Autocomplete: Triggers automatically as you type. Press `Tab` to accept.

- **In-place Edit (`Cmd+I`):** Select code, press `Cmd+I` to refactor selection inline.

- **Sidebar Chat (`Cmd+L`):** Ask questions, reference files with `@file` context.

SLIDE 5

Step 4: Ignore Indexes & Solve Latency

1. `.continueignore` Exclusions

Prevent 100% CPU loops by excluding dependencies from embeddings index:

```
node_modules/  
dist/  
build/  
.git/
```

2. Latency Prefix Tuning To solve autocomplete freeze when switching tabs, set:

- **Prefix Length:** 500
- **Suffix Length:** 250

SLIDE 6

Step 5: Terminal CLI Agents & Pipes



```
# 1. Launch terminal coding agents directly via Ollama
$ ollama launch opencode      # Autonomous developer TUI
$ ollama launch copilot-cli   # Command generator helper
$ ollama launch claude        # Agentic coding CLI tool

# 2. Or pipe log dumps directly for troubleshooting
$ cat error.log | \
  ollama run qwen2.5-coder:14b "Explain this error"
```

Ollama's desktop app includes a native registry to launch open-source terminal agents.

SLIDE 7

Quick-Reference Architecture Map

The Core Workflow

- **Tab Autocomplete:** Route to `Qwen-1.5B-Base` (runs at **188.4 tok/s**).
- **Chat Reasoning:** Route to `Qwen-14B-Instruct` (runs at **30.0 tok/s**).
- **CLI Pipe:** `cat error.log | ollama run qwen2.5-coder:14b`

Hardware Sizing & CPU Fixes

- **VRAM Rule:** Keep models within VRAM bounds to prevent swap lags.
- **VS Code Settings:** Set autocomplete Prefix to `500` & Suffix to `250`.
- **CPU Protection:** Prevent CPU loops by adding compiled folders to `.continueignore`.

SLIDE 8

Reclaim Your Code Privacy

Full configurations, custom `ModelFile` setups, and terminal pipes are in the full guide.

Connect & Build:

- **Substack:** softwarepermanence.substack.com
- **GitHub (Companion Code):**
github.com/praveenveera/software-permanence
- **Dev.to:** dev.to/praveen_builds
- **Medium:** medium.com/@praveenveera92
- **Instagram:** [@praveen_builds](https://www.instagram.com/praveen_builds)
- **LinkedIn:** [linkedin.com/in/praveen-veera-6ab22567](https://www.linkedin.com/in/praveen-veera-6ab22567)
- **Hashnode:** hashnode.com/@praveen-builds